Since 2004
UET
ĐẠI HỌC CÔNG NGHỆ, ĐHQGHN
VNU-University of Engineering and Technology

Since 1906
VNU
ĐẠI HỌC QUỐC GIA HÀ NỘI
Vietnam National University, Hanoi

INT3121 – Topics in Computer Science
Lecture 1: Foundation of LLMs

Hanoi, 09/2025

Course Introduction

- Introduction to Large Language Models (LLMs)
- Word embedding
- Tokenization
- Self-attention
- Muli-headed self-attention
- Vision Transformer

From AI to Large Language Models

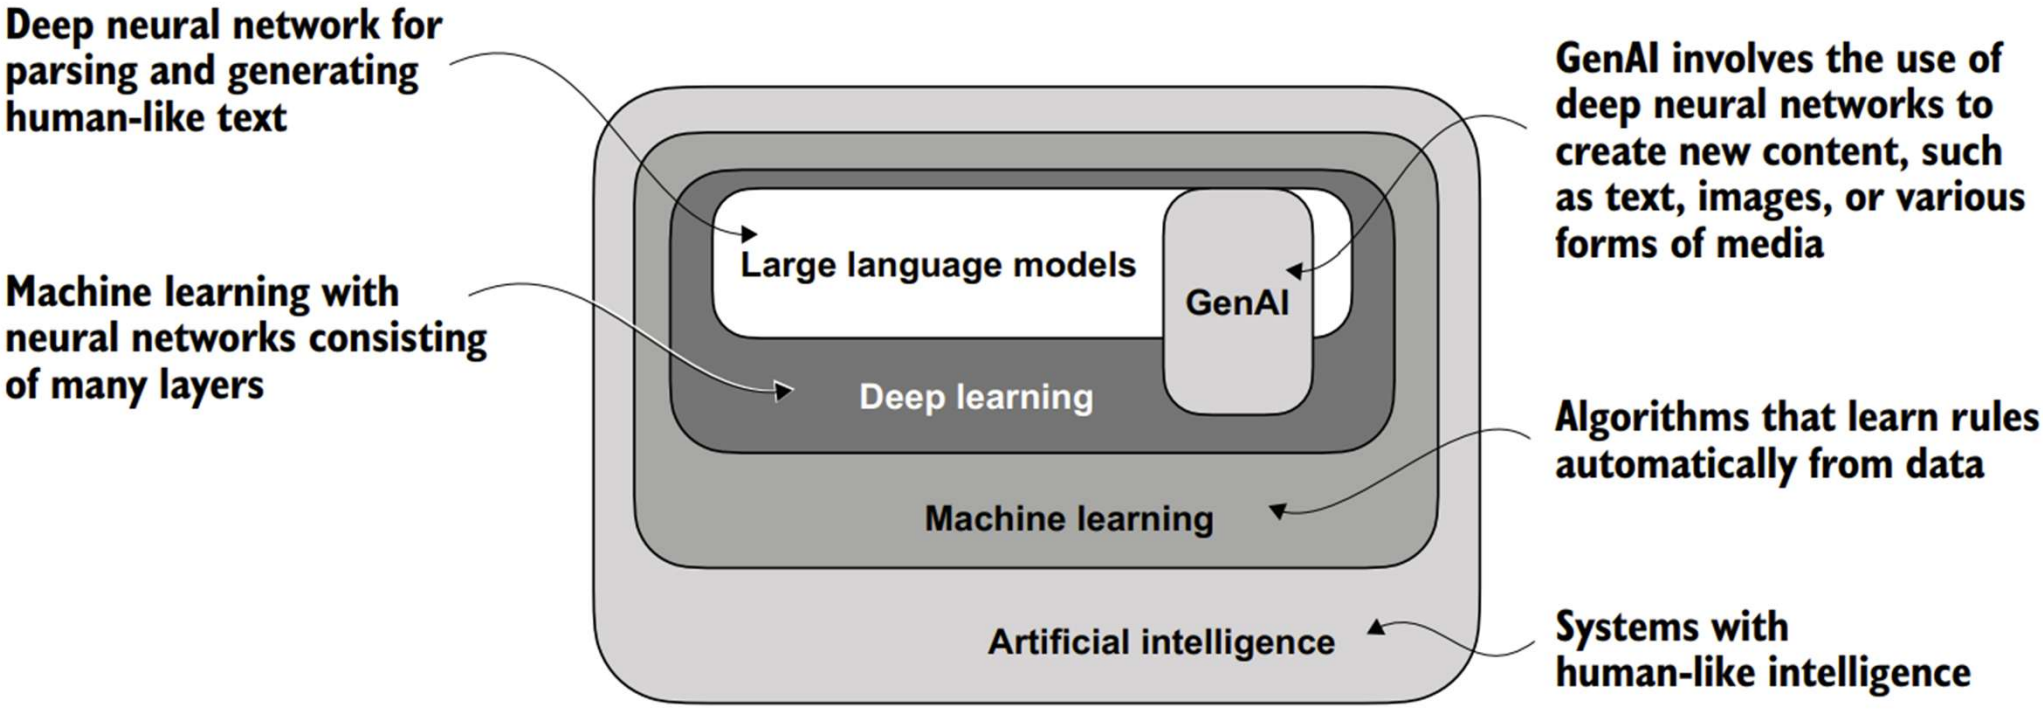

Source: BLLM

Source: COS 597G

Application of LLMs

- Chatbot (ChatGPT, Gemini,…)
- Machine translation
- Generation of novel texts
- Sentiment analysis
- Text summarization
- Language-instructed image generation
- Domain-specific LLM-based applications
  - Code generation (Copilot)
  - Task/project magagement (Notion)
  - Multimedia generation/editing (Canva)

Source: CCS-597G

Language Models: Narrow Sense

- A probabilistic model that assigns a probability $P[w_1, w_2, \ldots, w_n]$ to every finite sequence $w_1, \ldots, w_n$ (grammatical or not)

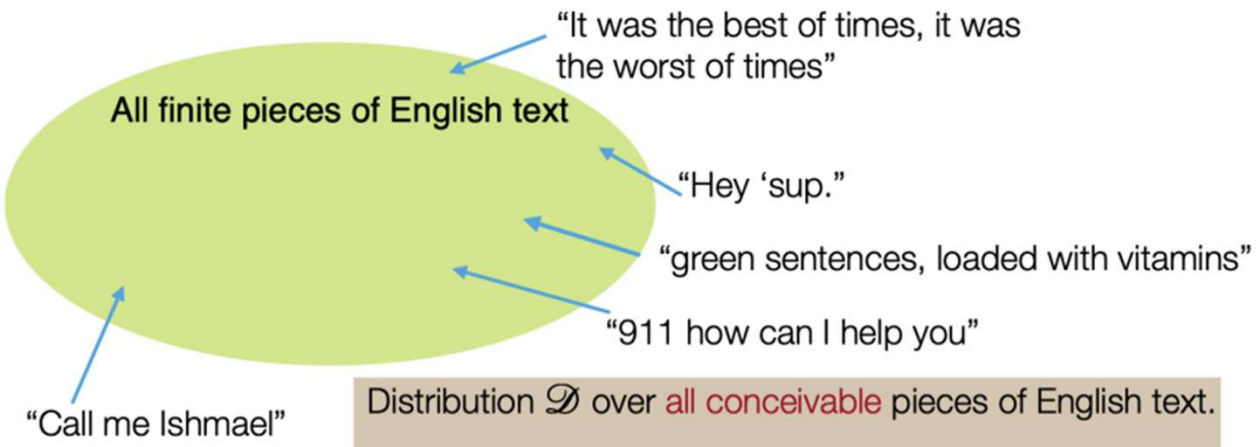

Distribution $\mathscr{D}$ over all conceivable pieces of English text.

Assigns probability $\mathbf{Pr[w_1 w_2 w_2 \ldots w_n]}$ to every finite word sequence $\mathbf{w_1 w_2 w_2 \ldots w_n}$ (grammatical or not).

Source: COS 324

Language Models: Narrow Sense

Conditional probability

$$p(w_1, w_2, w_3, \ldots, w_N) =$$
$$p(w_1)\, p(w_2|w_1)\, p(w_3|w_1, w_2) \times \ldots \times p(w_N|w_1, w_2, \ldots w_{N-1})$$

Sentence: “the cat sat on the mat”

$$P(\text{the cat sat on the mat}) = P(\text{the}) * P(\text{cat}|\text{the}) * P(\text{sat}|\text{the cat})$$
$$* P(\text{on}|\text{the cat sat}) * P(\text{the}|\text{the cat sat on})$$
$$* P(\text{mat}|\text{the cat sat on the})$$

Implicit order

Source: COS 484

GPT-3 still acts in this way but the model is implemented as a very large neural network of 175-billion parameters!

The Transformer Architecture: Overview

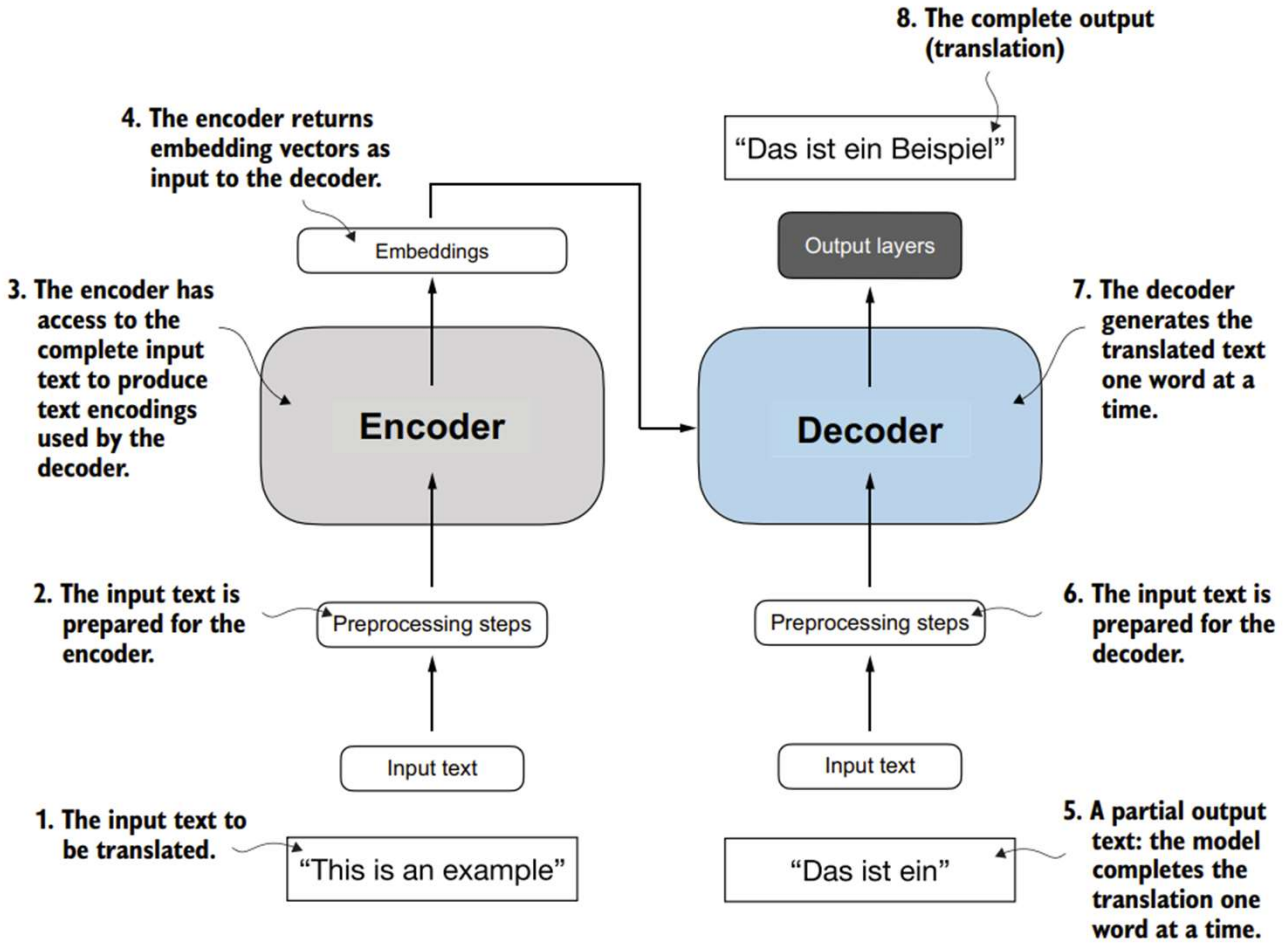

The Transformer Architecture: Detail

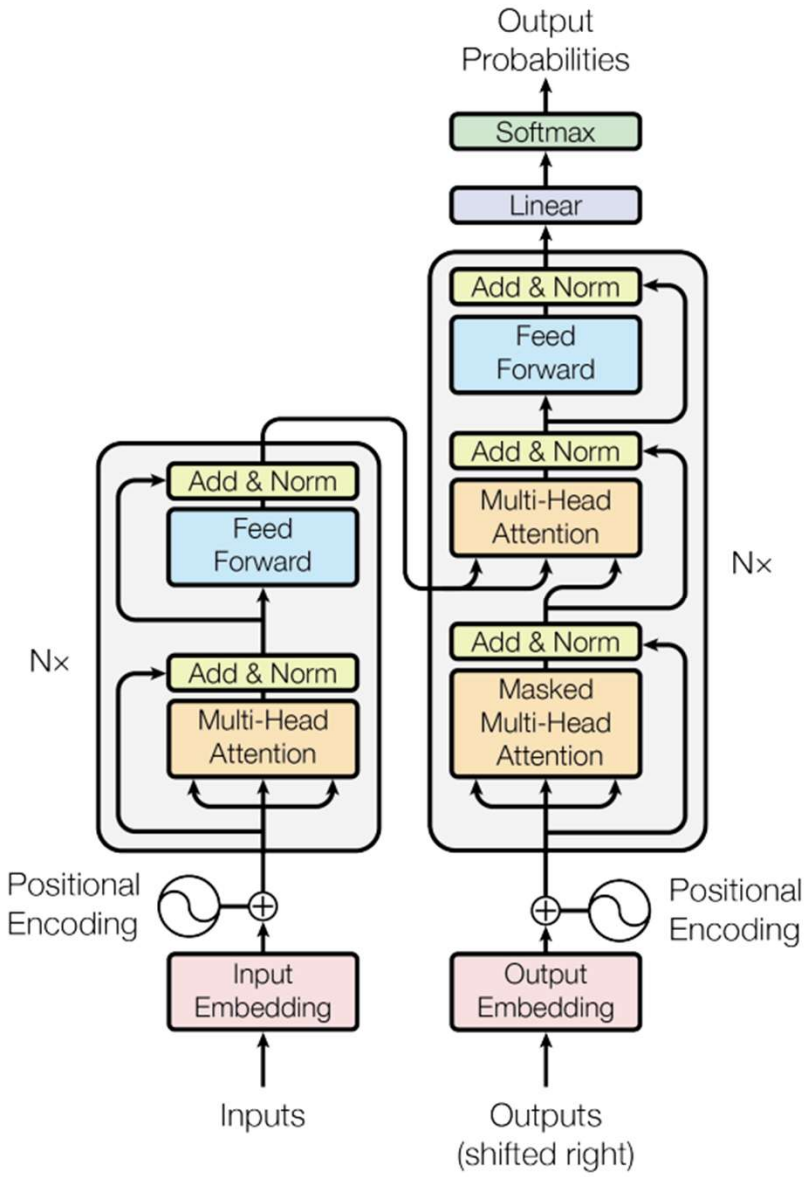

Language Models: Broad Sense

- Decoder-only models (GPT-x models)
- Encoder-only models (BERT, RoBERTa, ELECTRA)
- Encoder-decoder models (T5, BART)

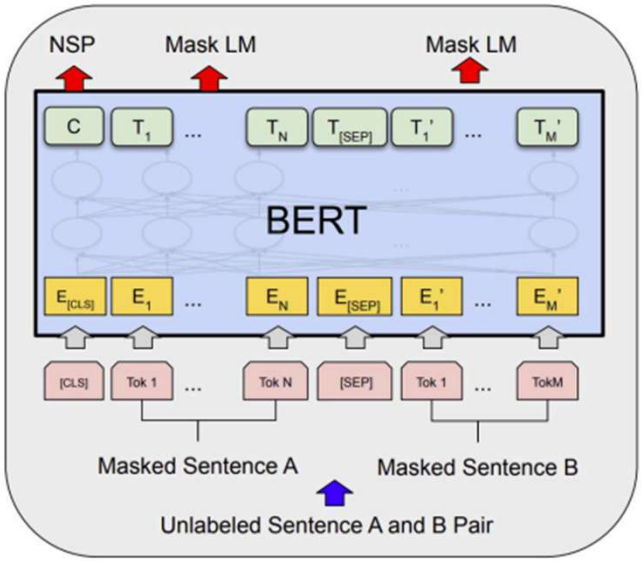

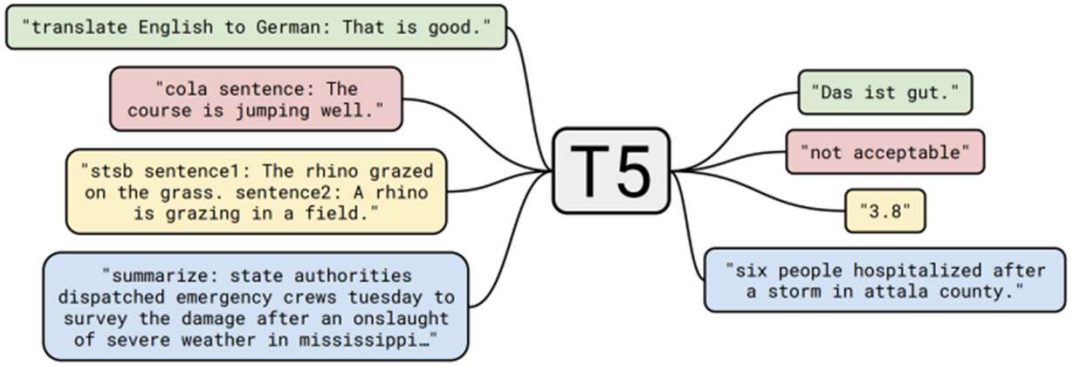

Source: COS 597G

GPT vs BERT

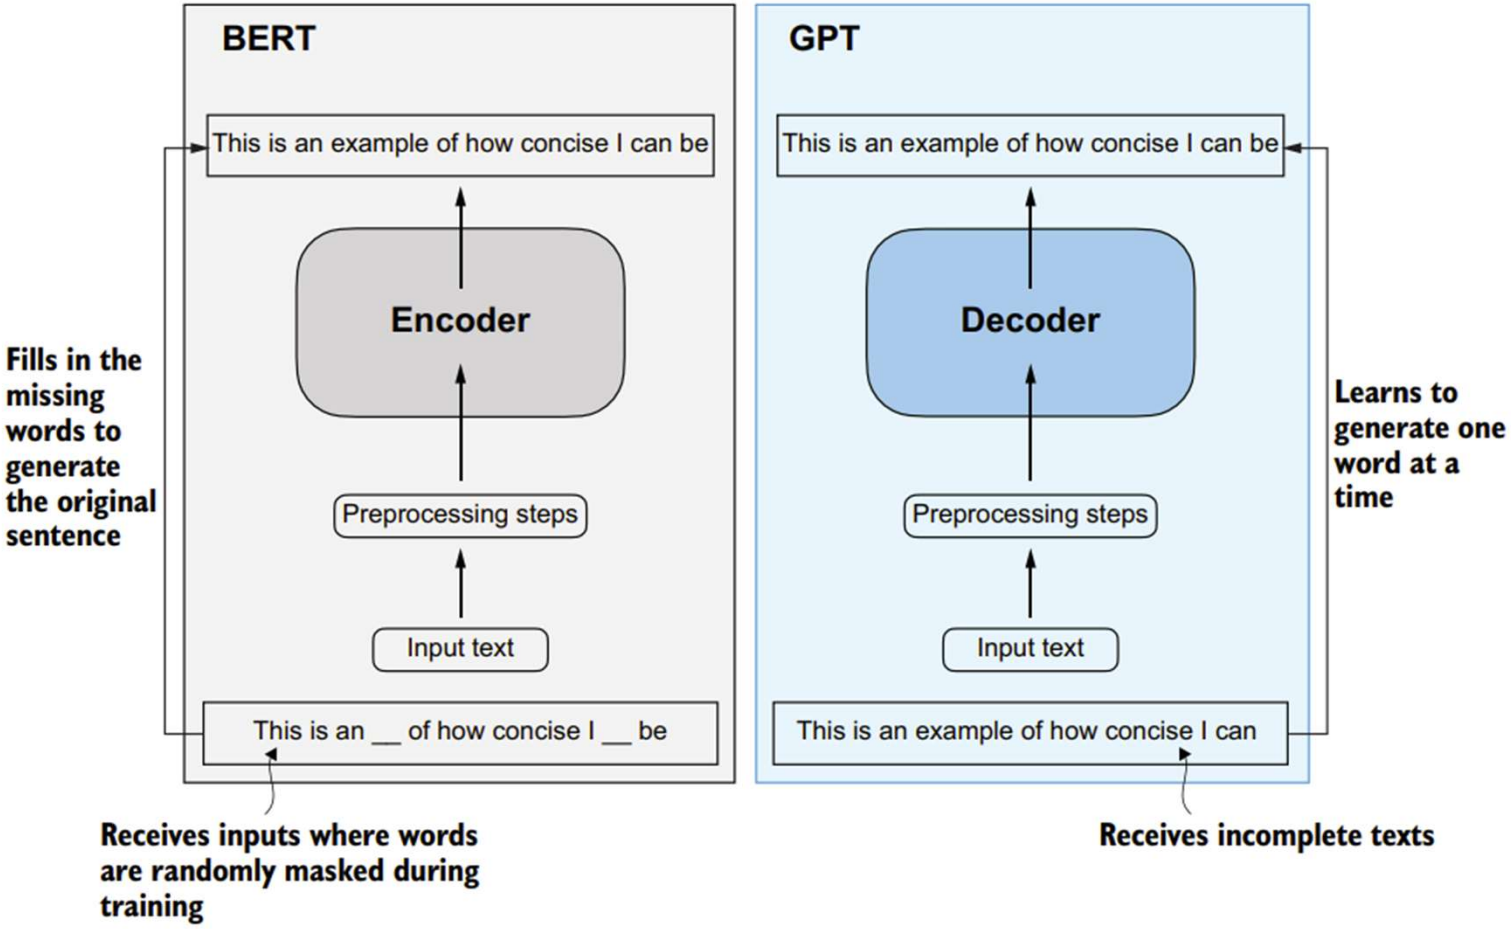
BERT
This is an example of how concise I can be
Encoder
Preprocessing steps
Input text
This is an __ of how concise I __ be
Fills in the missing words to generate the original sentence
Receives inputs where words are randomly masked during training
GPT
This is an example of how concise I can be
Decoder
Preprocessing steps
Input text
This is an example of how concise I can
Learns to generate one word at a time
Receives incomplete texts

GPT-based Model Capabilities

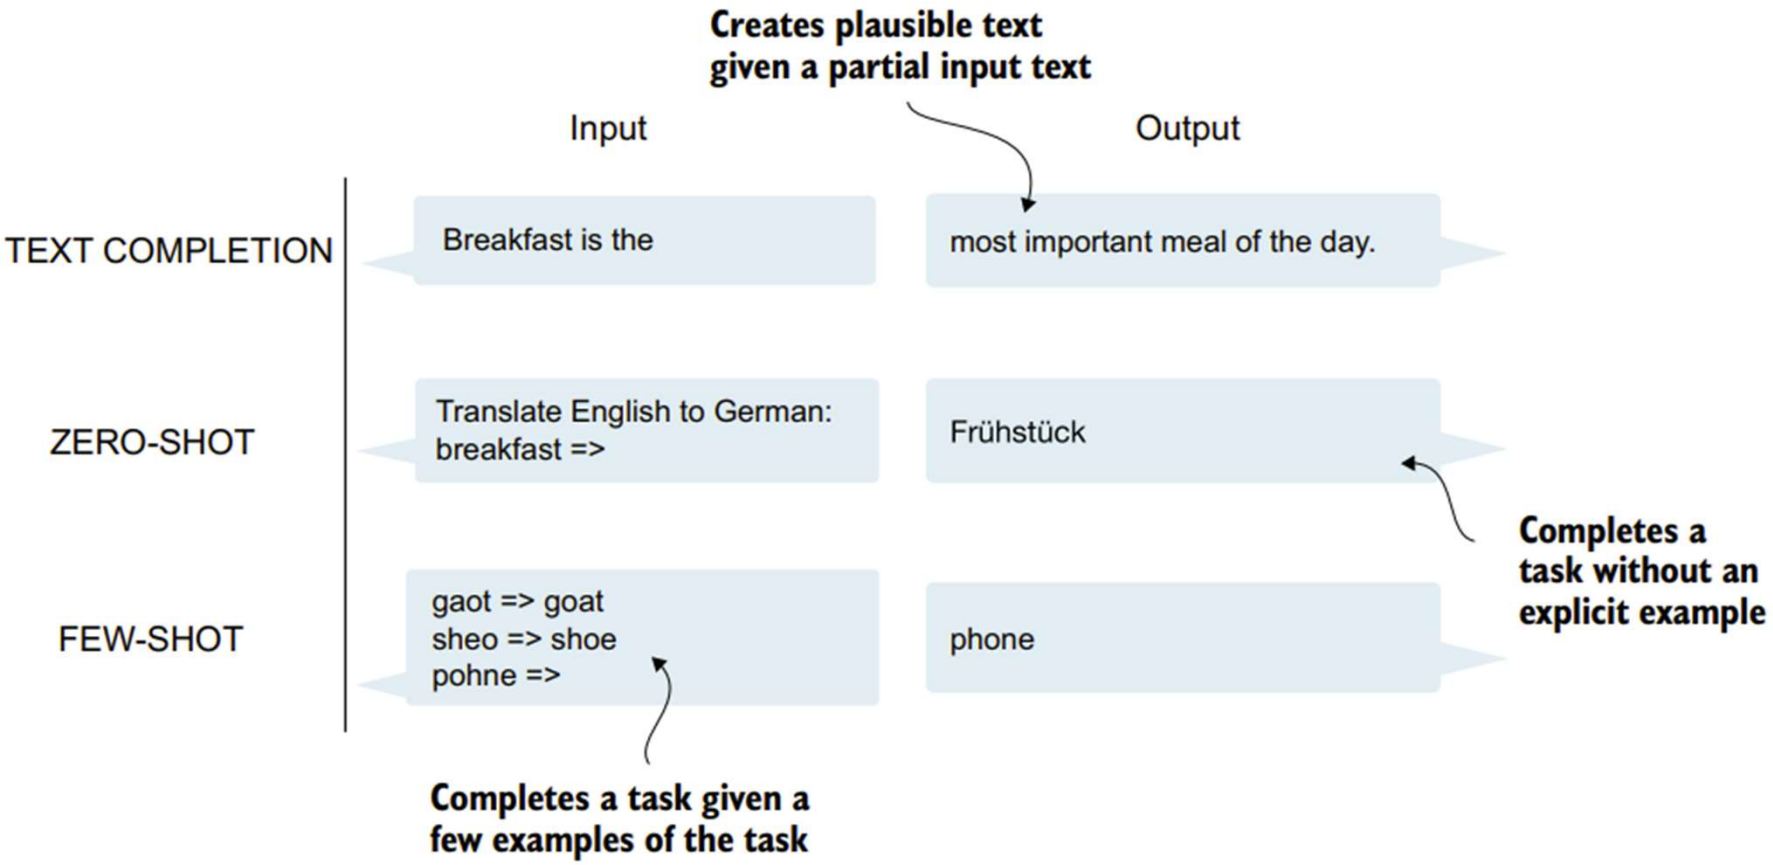

Stages of Building and Using LLMs

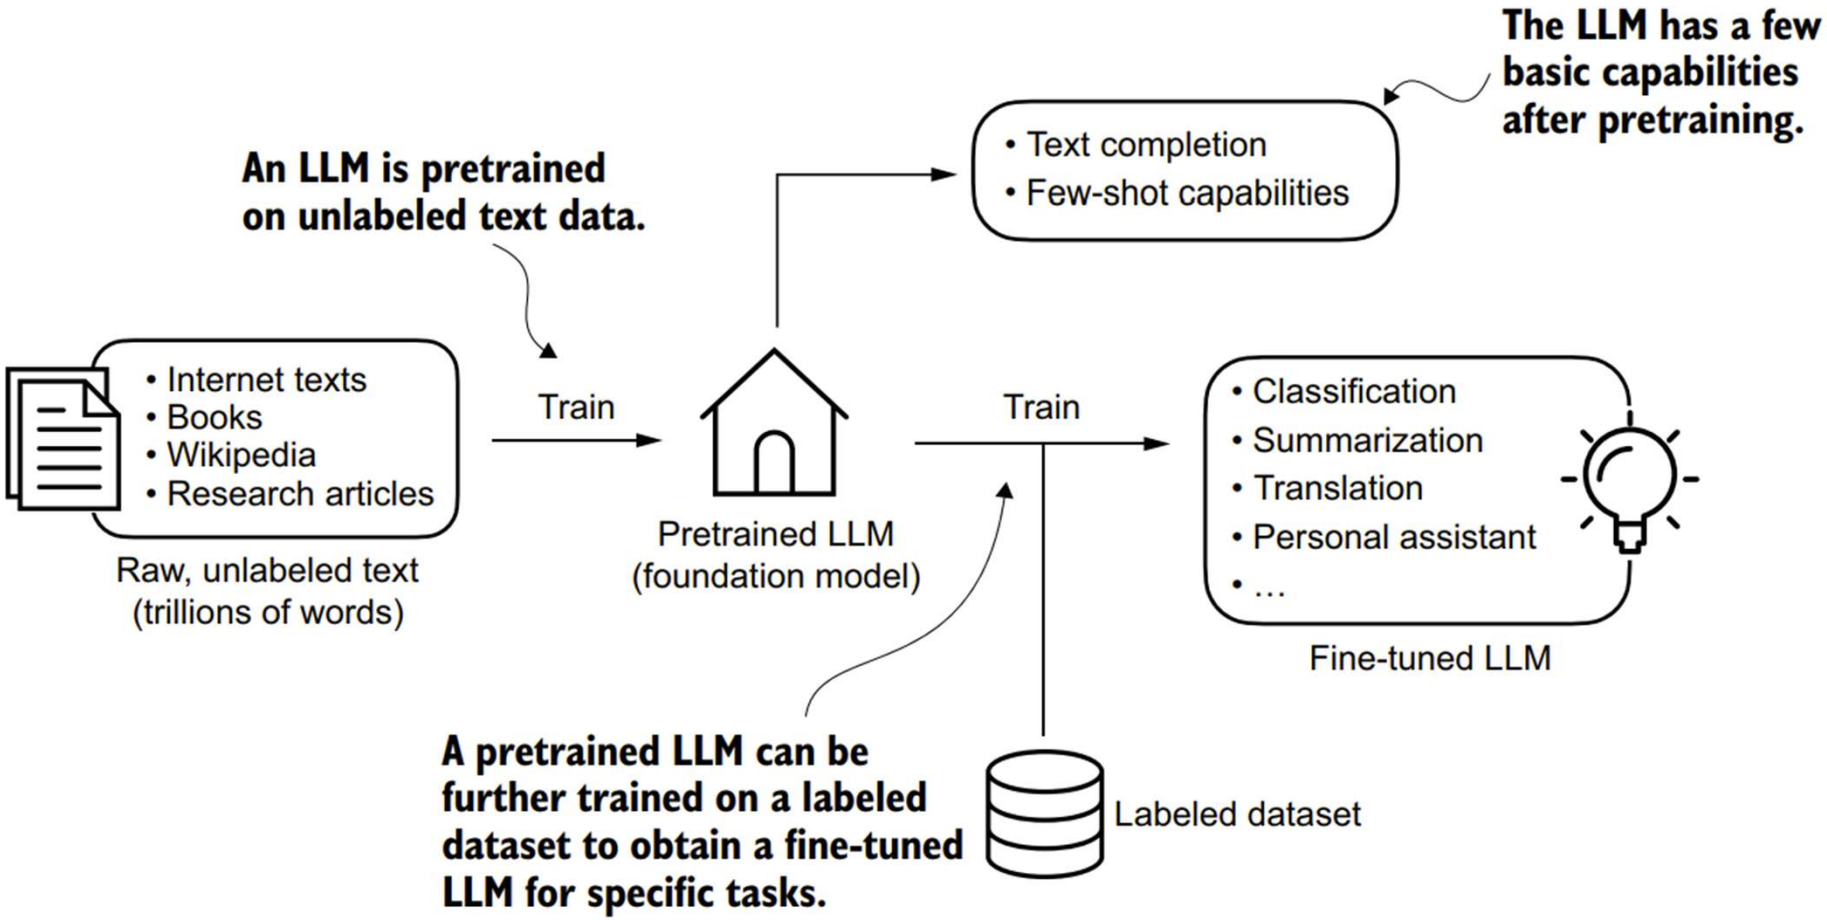

GPT 3 Training Dataset

Table 1.1 The pretraining dataset of the popular GPT-3 LLM

Dataset name	Dataset description	Number of tokens	Proportion in training data
CommonCrawl (filtered)	Web crawl data	410 billion	60%
WebText2	Web crawl data	19 billion	22%
Books1	Internet-based book corpus	12 billion	8%
Books2	Internet-based book corpus	55 billion	8%
Wikipedia	High-quality text	3 billion	3%

GPT-3 has 175 billion parameters

LLMs Comparison

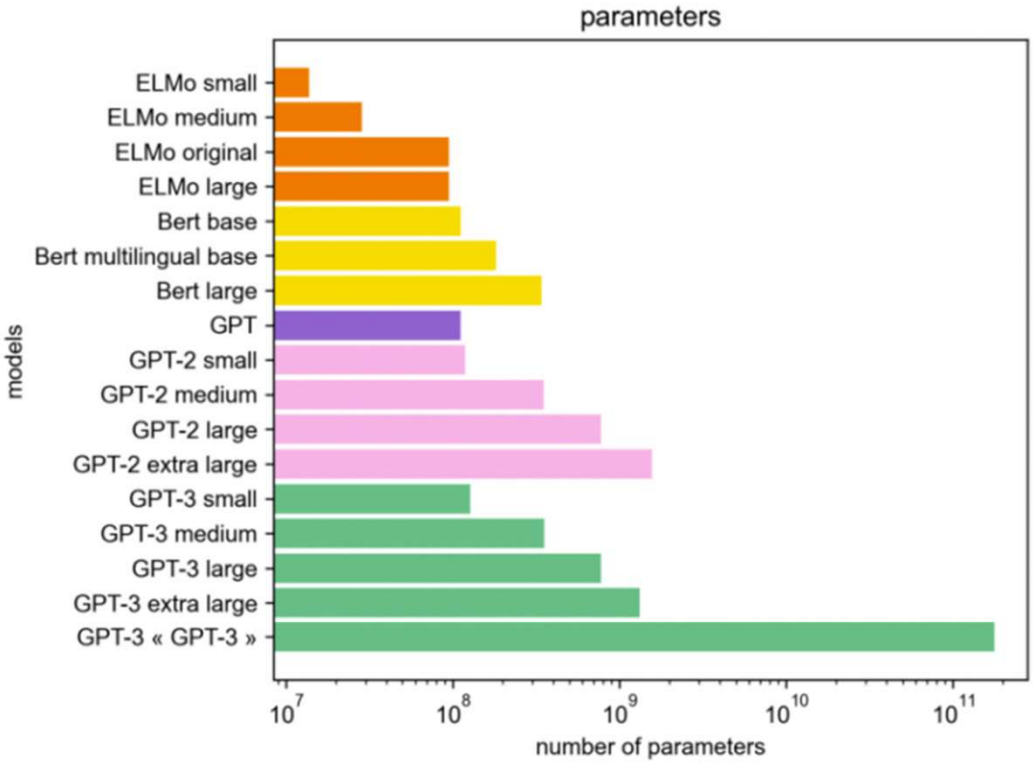

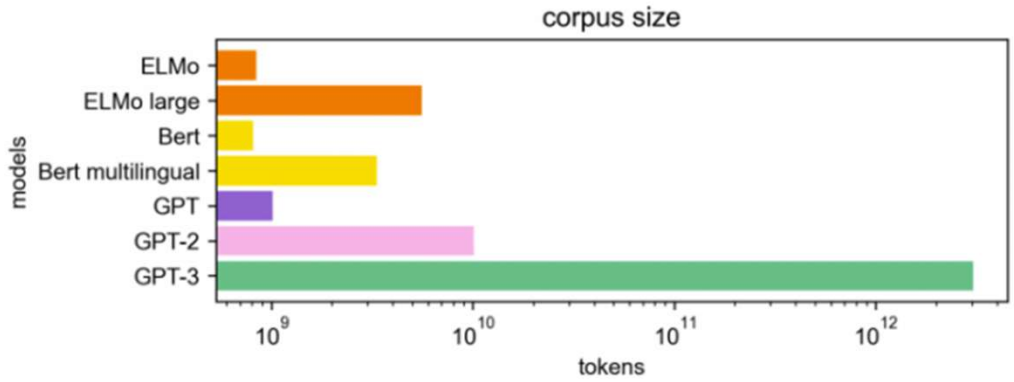

More recent models: PaLM (540B), OPT (175B), BLOOM (176B)…

Image source: https://hellofuture.orange.com/en/the-gpt-3-language-model-revolution-or-evolution/

A Closer Look at the GPT Architecture

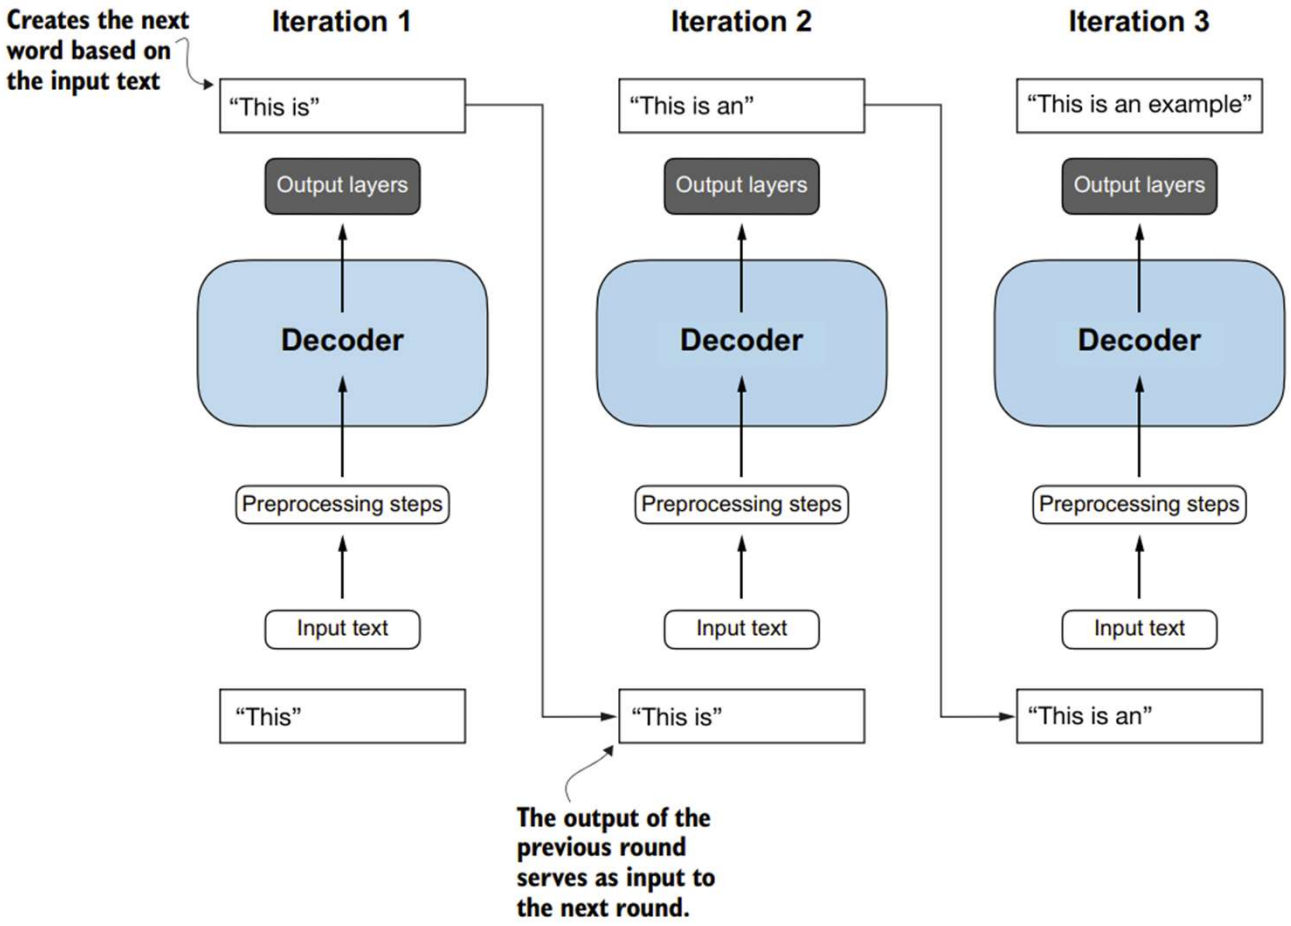

Implementing LLMs

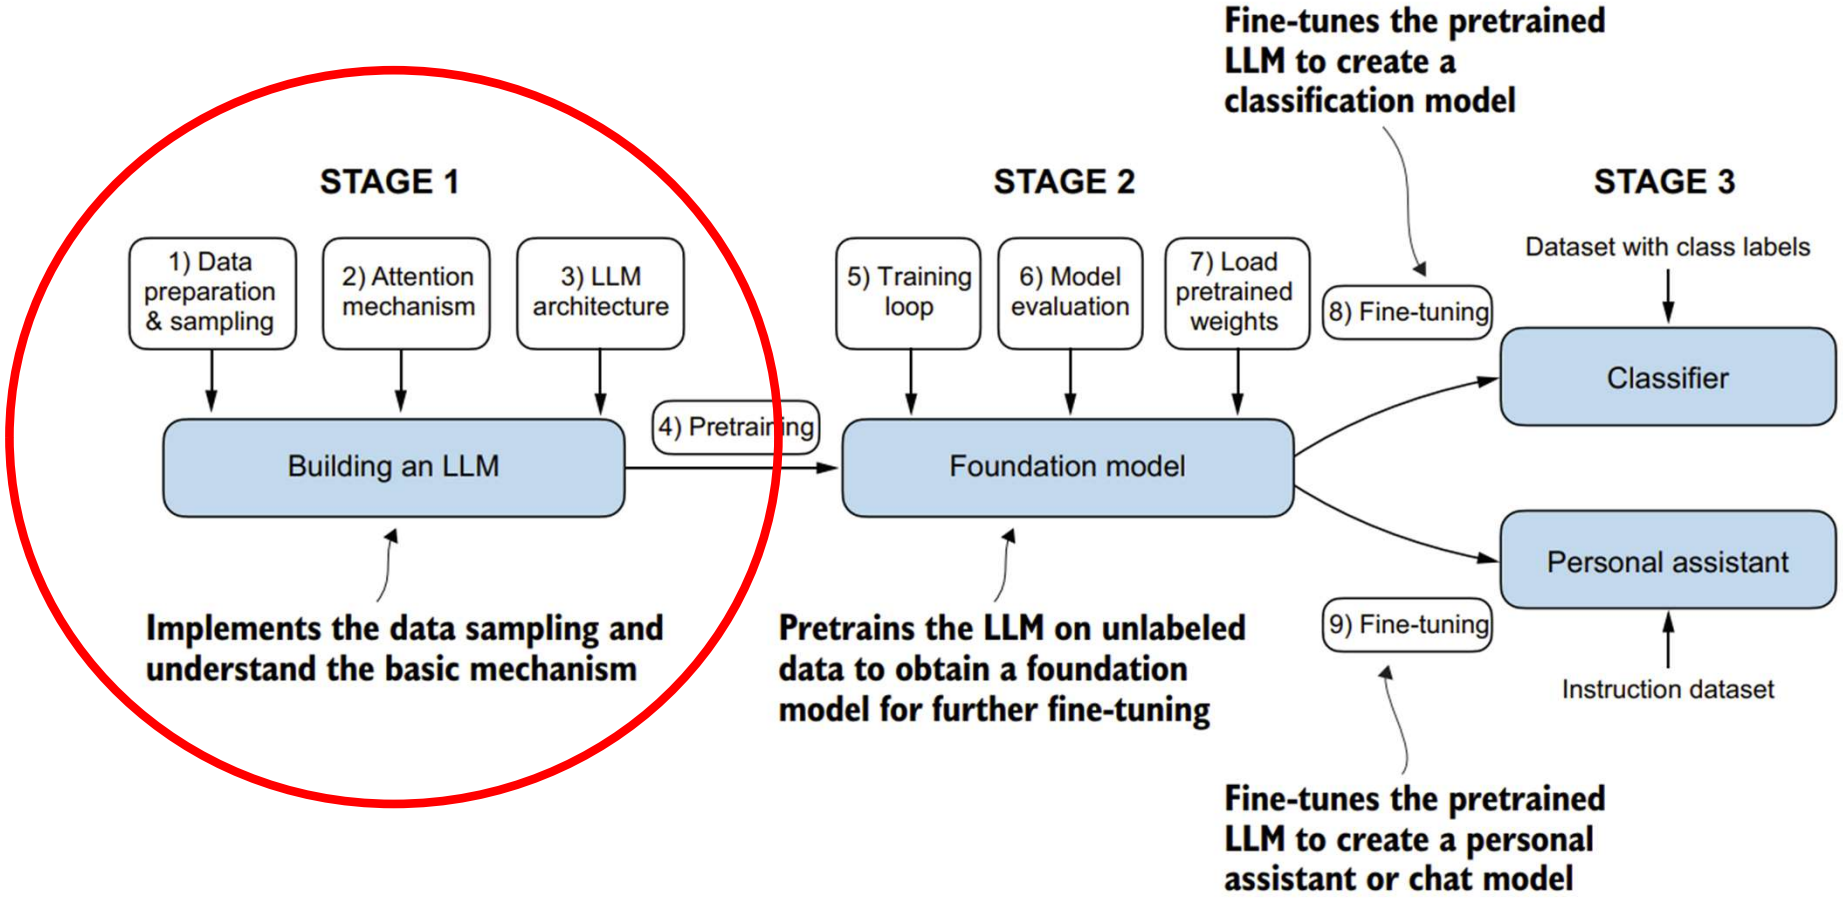

Word embedding – Modelling word in computer

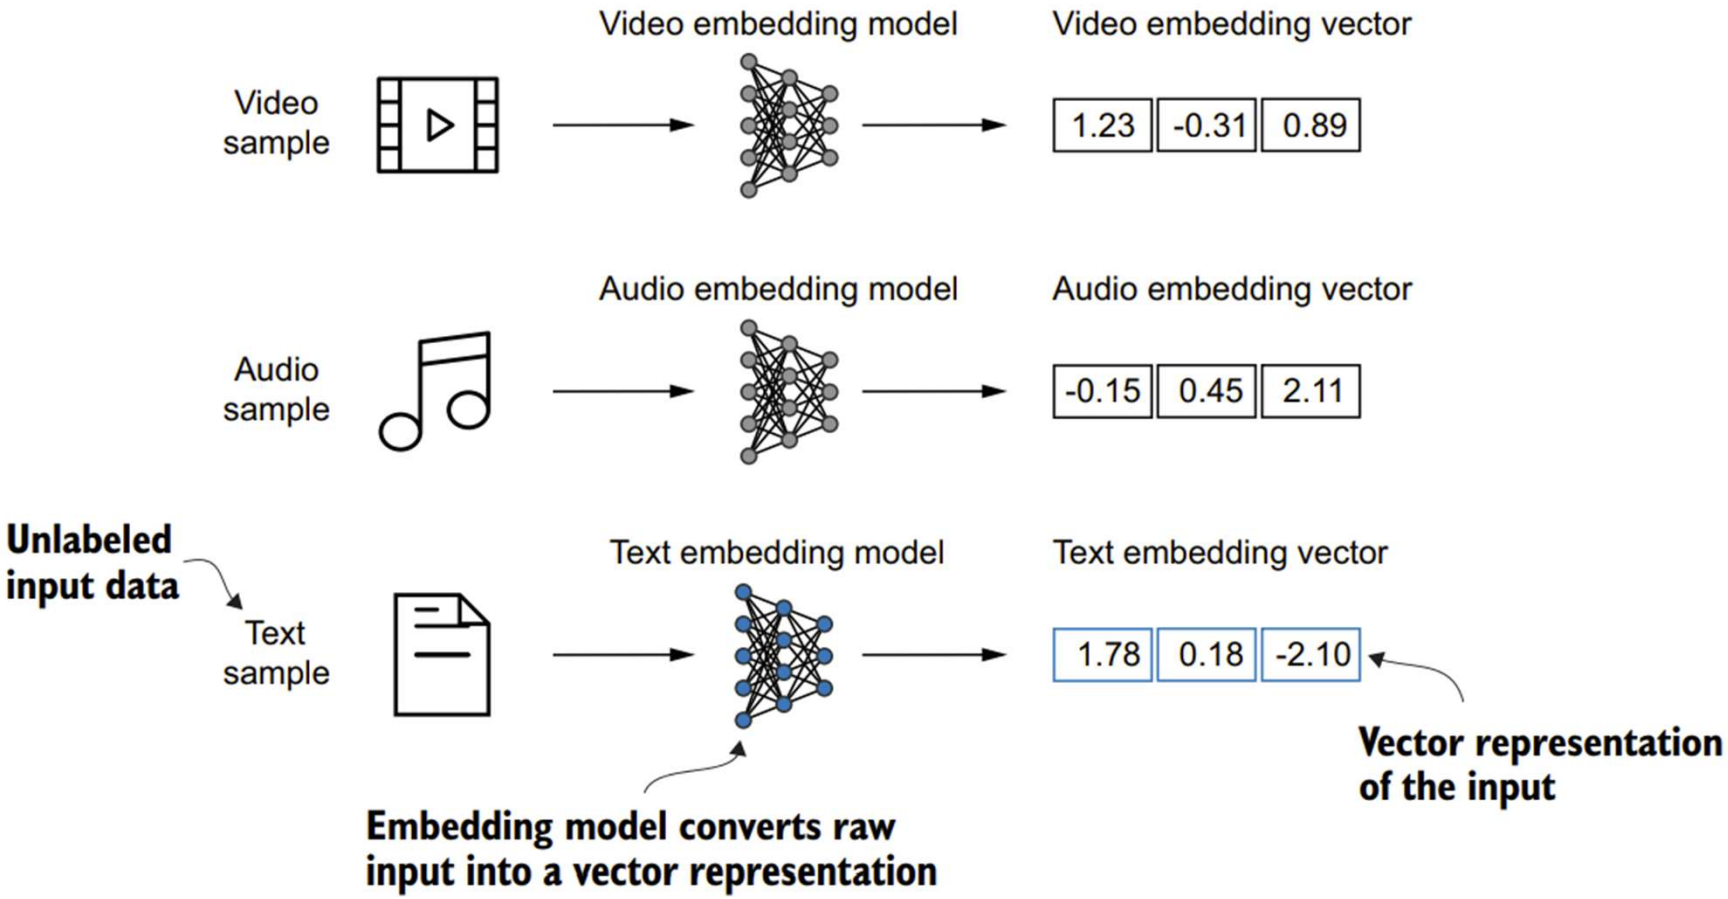

Word embedding – Modelling word in computer

Bringing words that have closer meaning closer in the embedding space

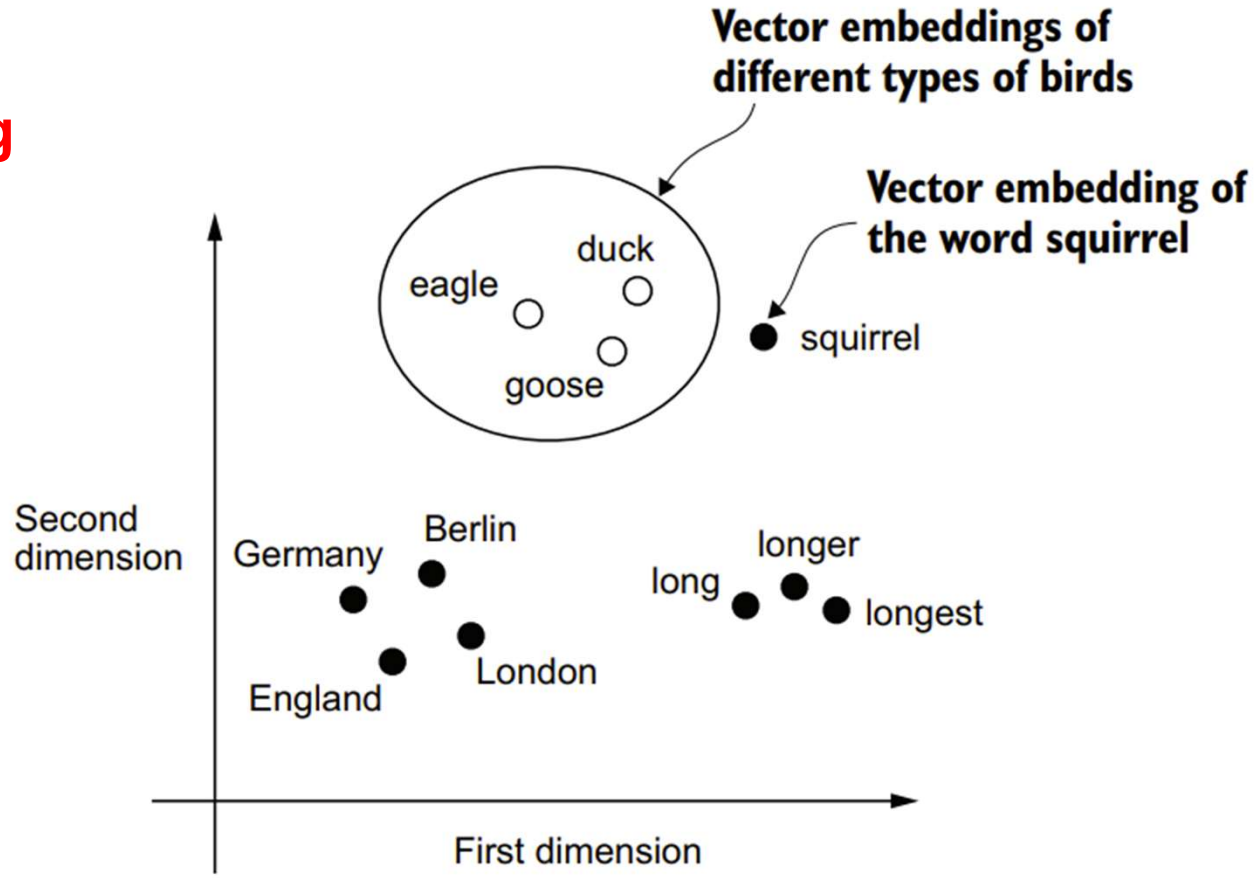

Tokenization

ChatGPT: Byte Pair Encoding (BPE)

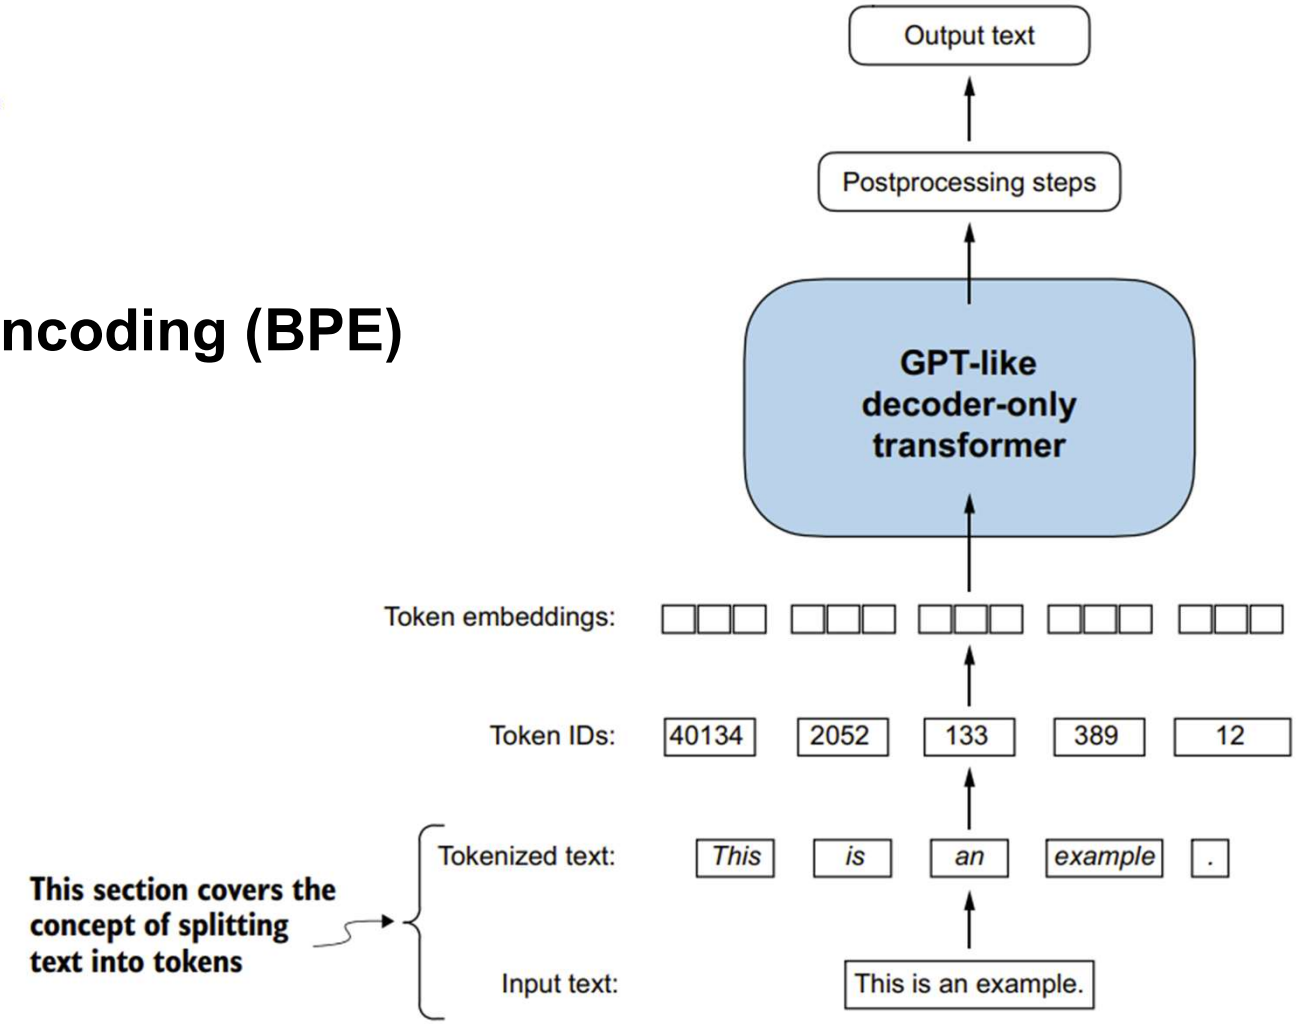

Tokenization in action

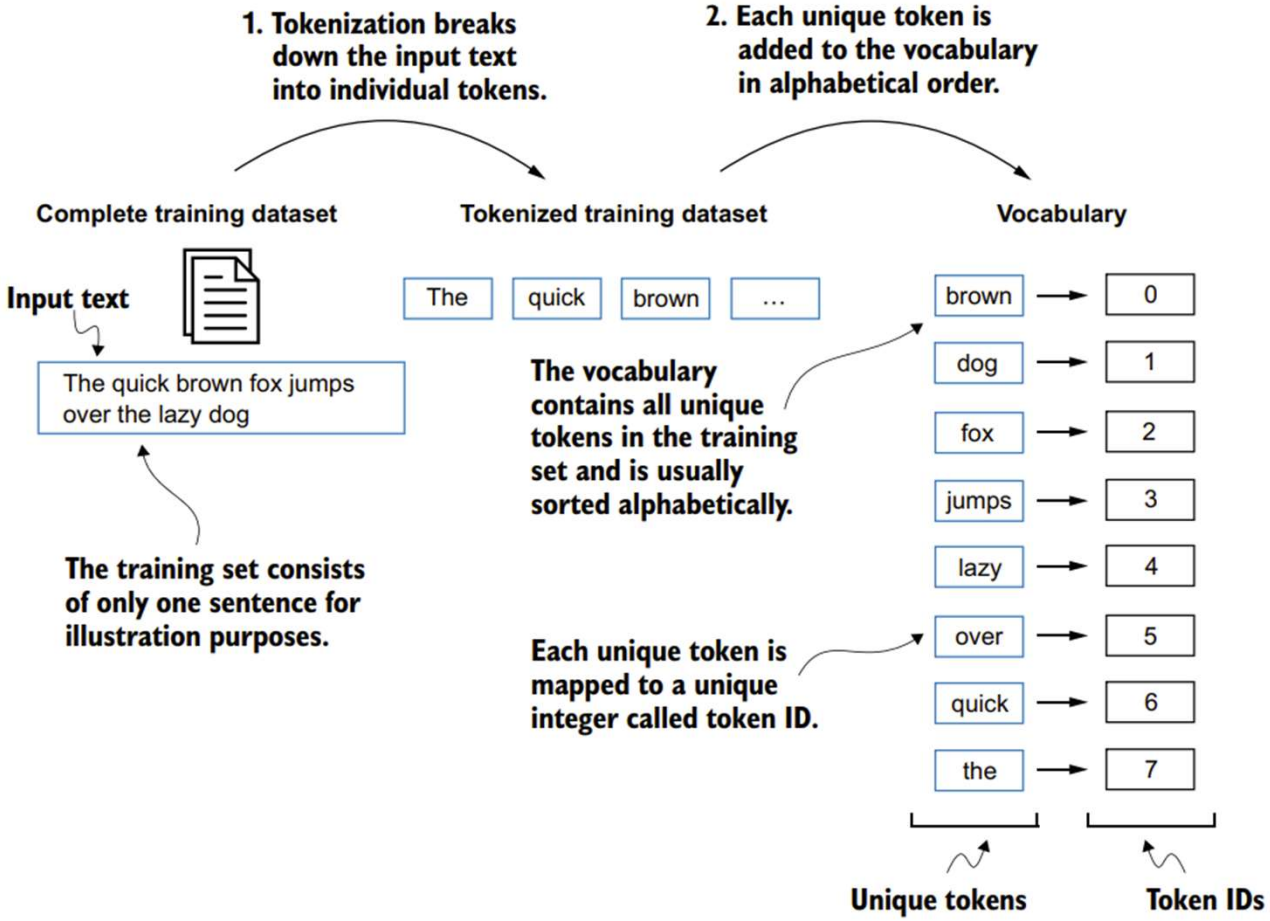

Input to LLM

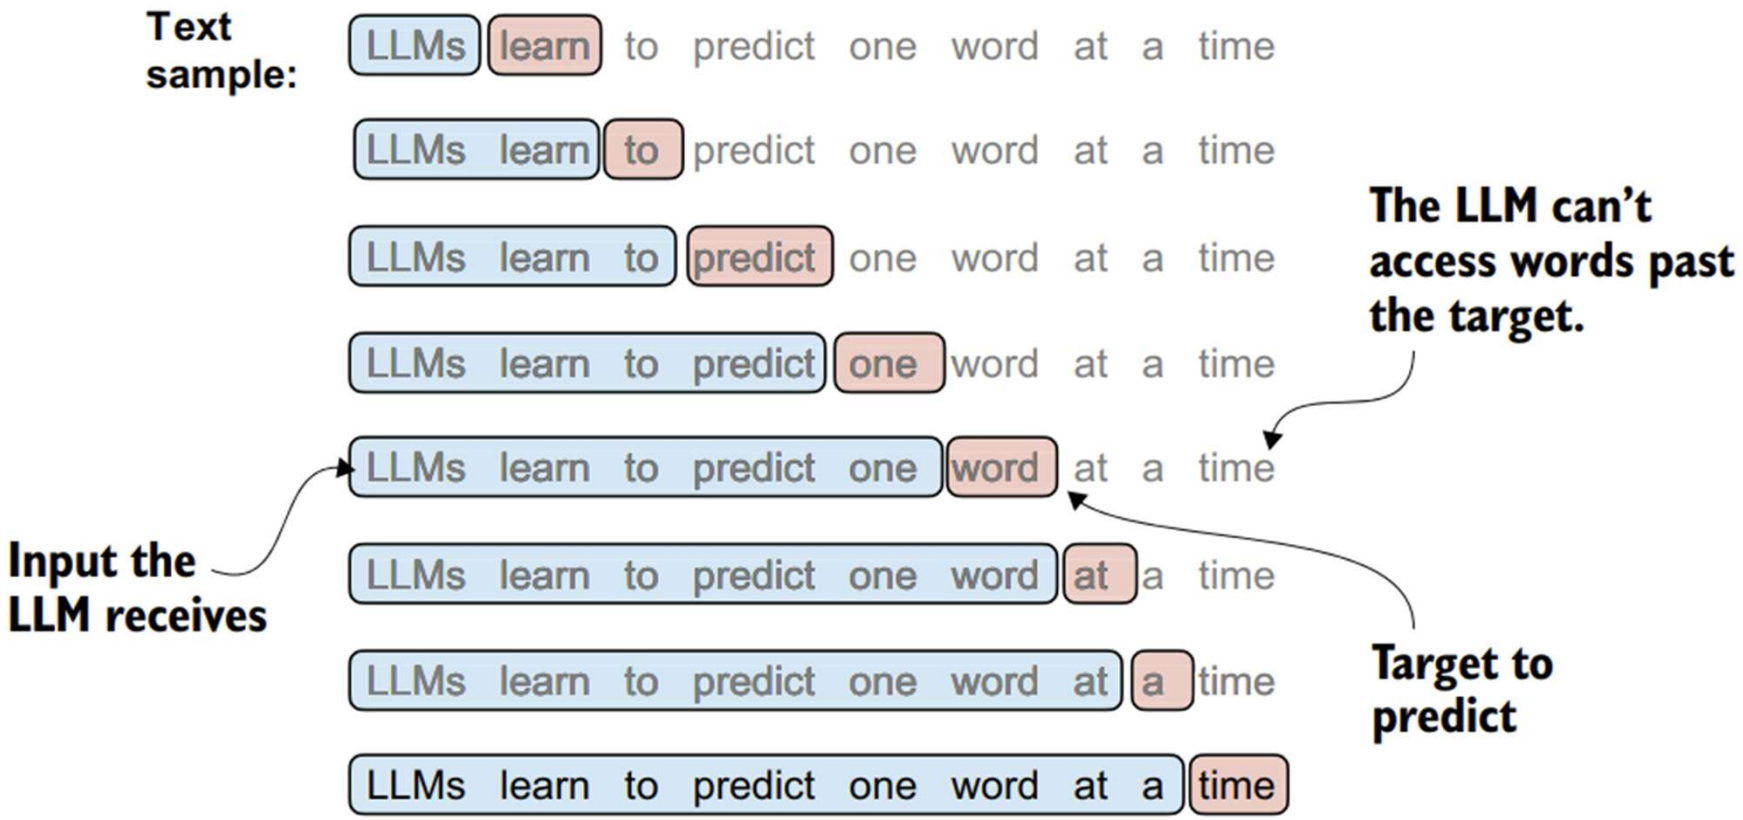
Text sample:
LLMs learn to predict one word at a time
LLMs learn to predict one word at a time
LLMs learn to predict one word at a time
LLMs learn to predict one word at a time
LLMs learn to predict one word at a time
LLMs learn to predict one word at a time
LLMs learn to predict one word at a time
LLMs learn to predict one word at a time
The LLM can't access words past the target.
Input the LLM receives
Target to predict

Self-Attention mechanism

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^\top}{\sqrt{d_k}})V$$

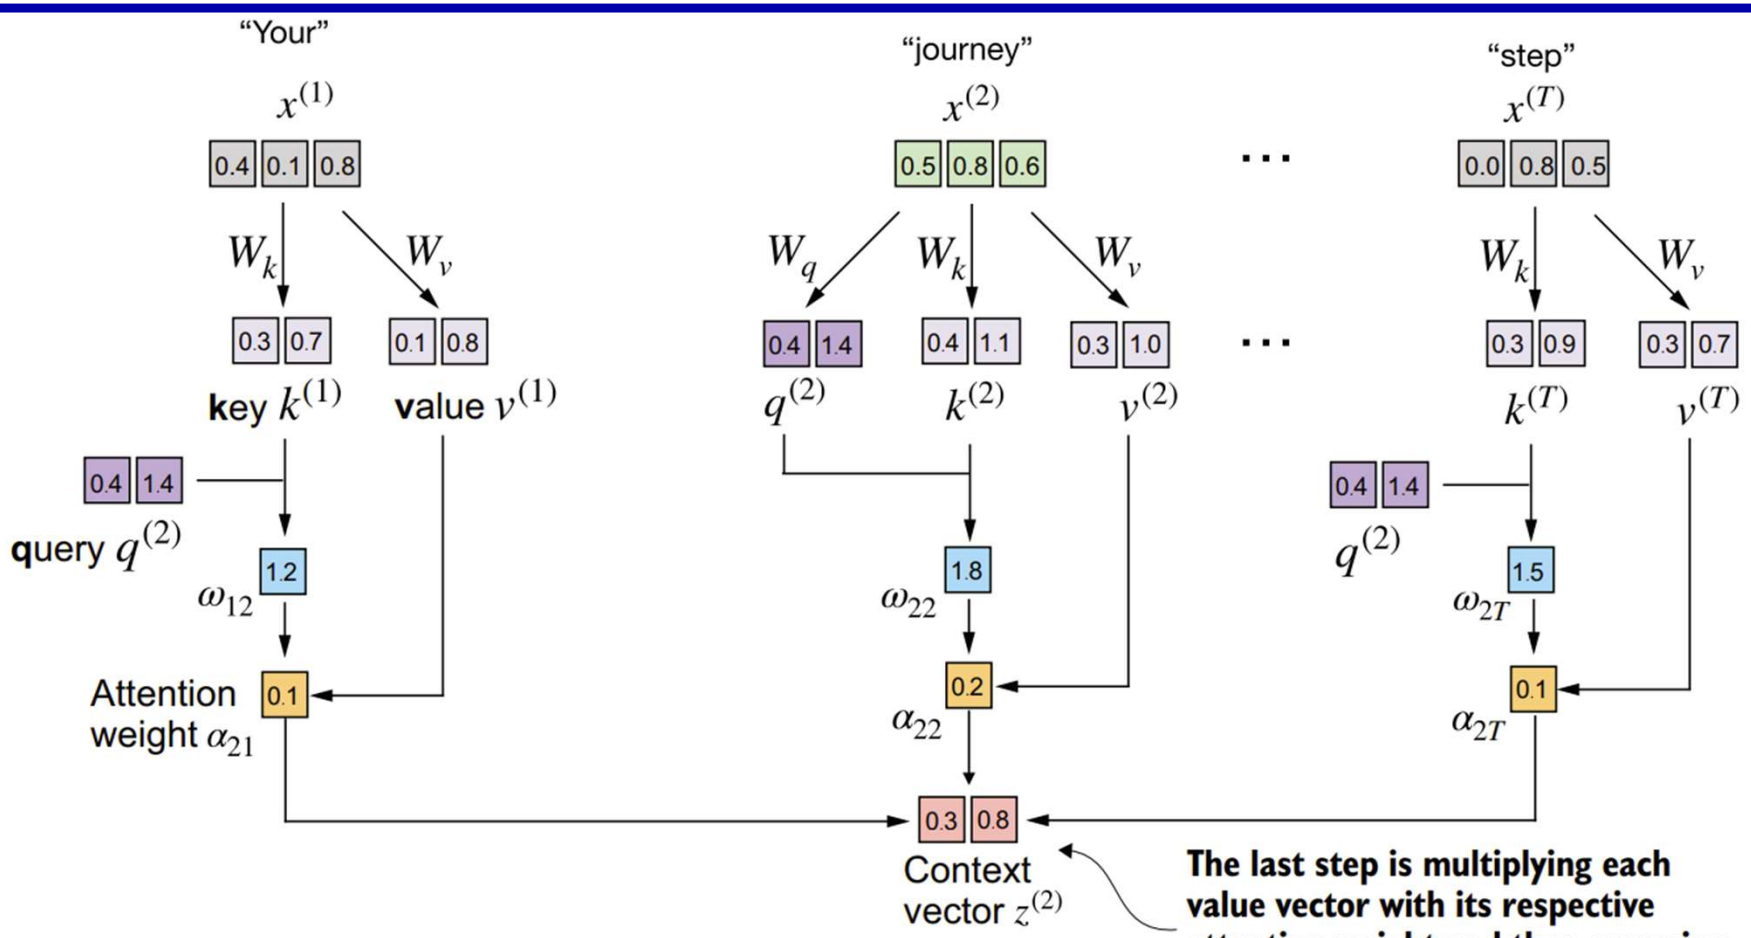

Self-Attention Overall Picture

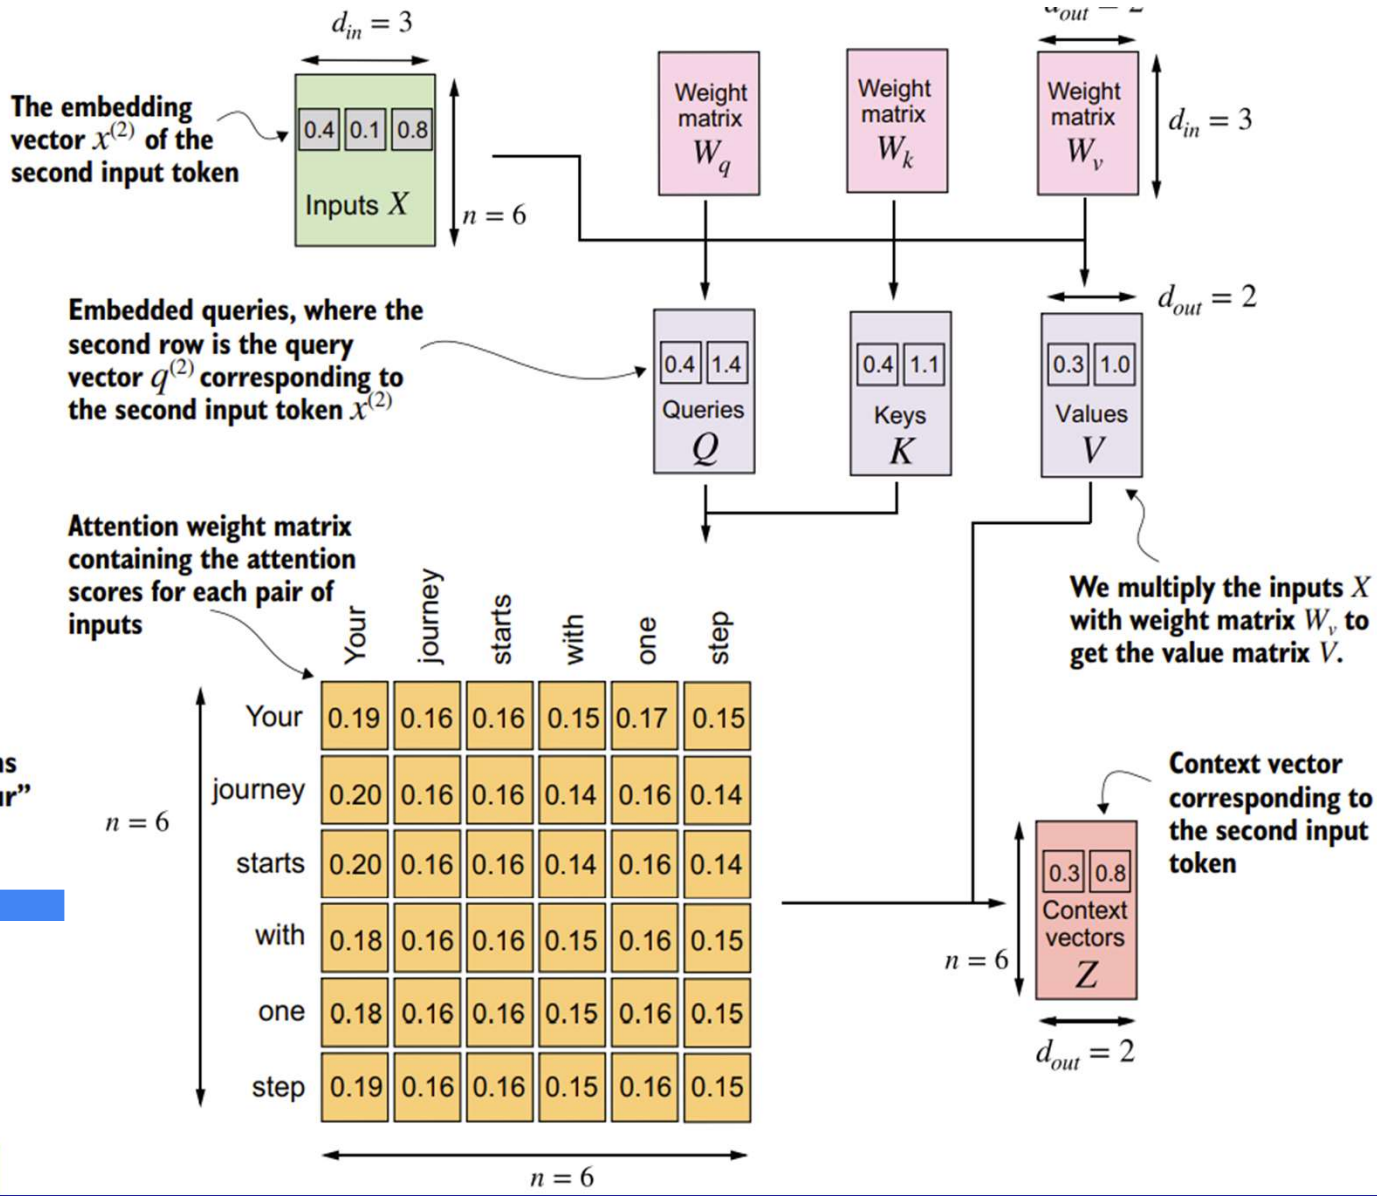

Causal attention

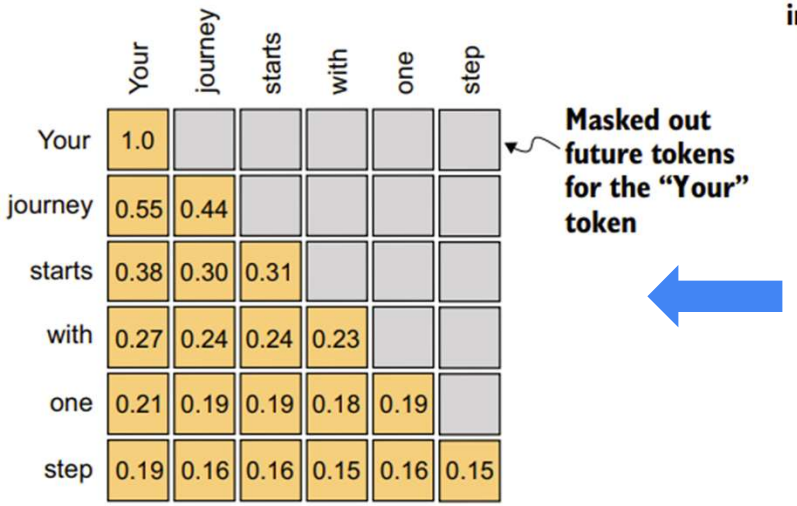

Self-Attention mechanism with dropout

Help prevent overfitting

Only used during training

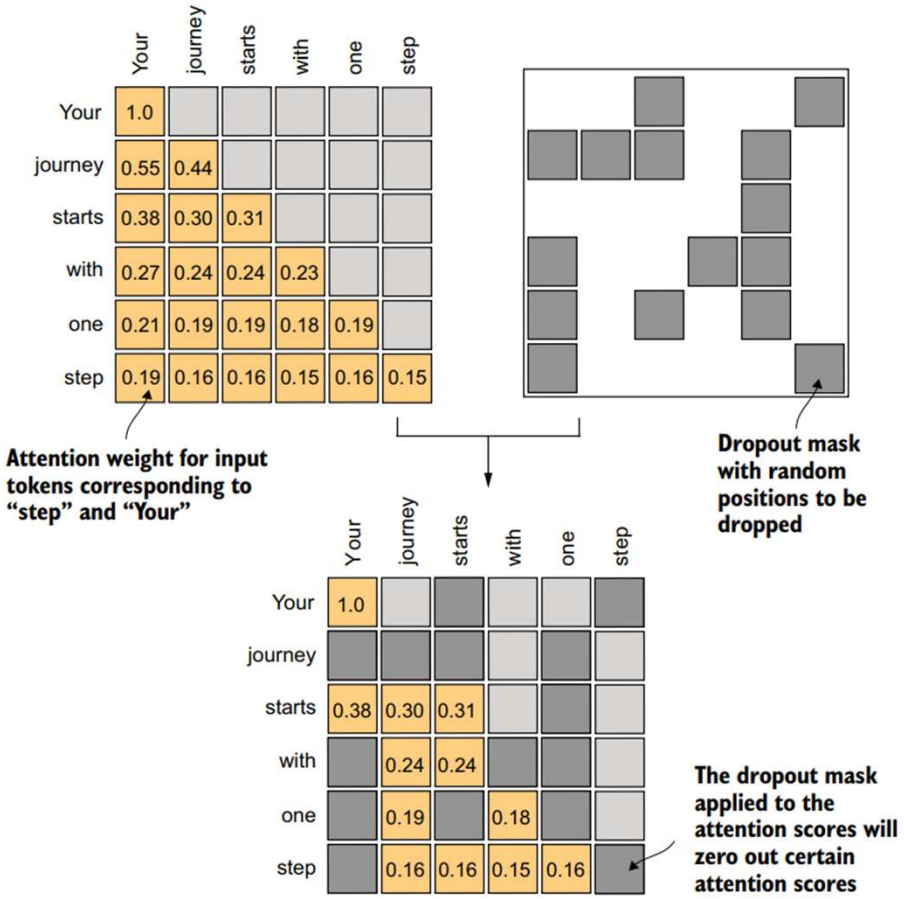

Multi-headed self-attention mechanism

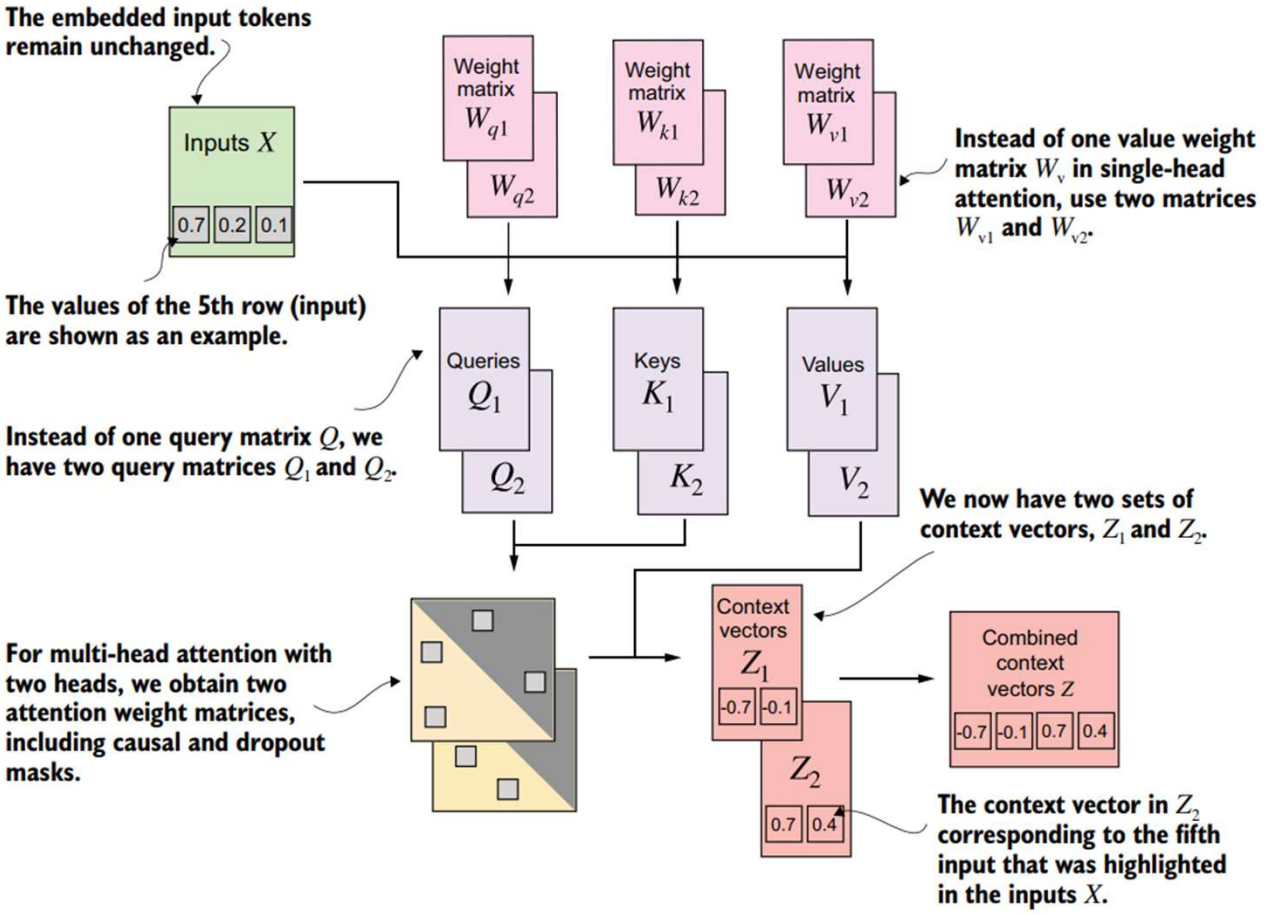

Multi-headed self-attention mechanism

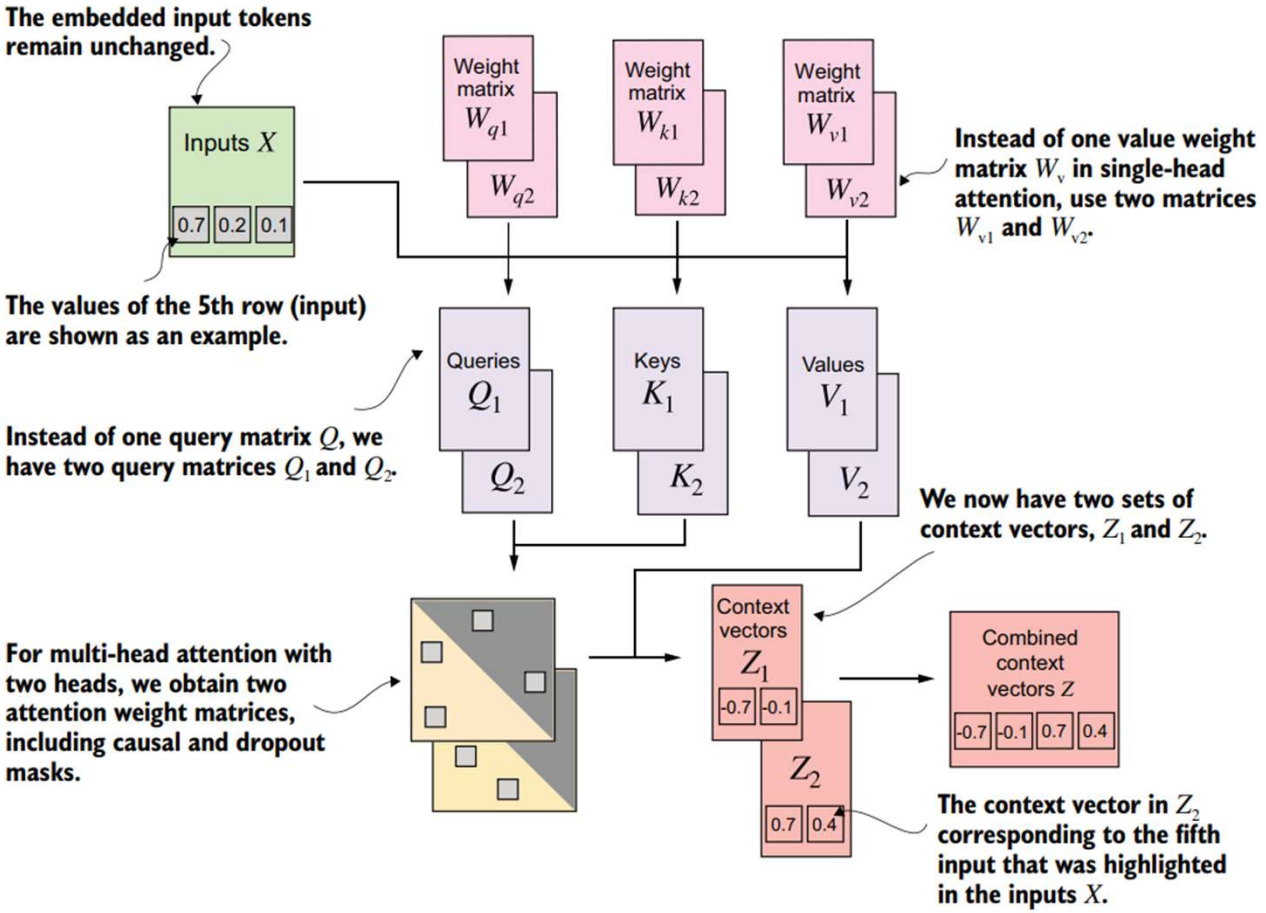

Vision Transformer (ViT)

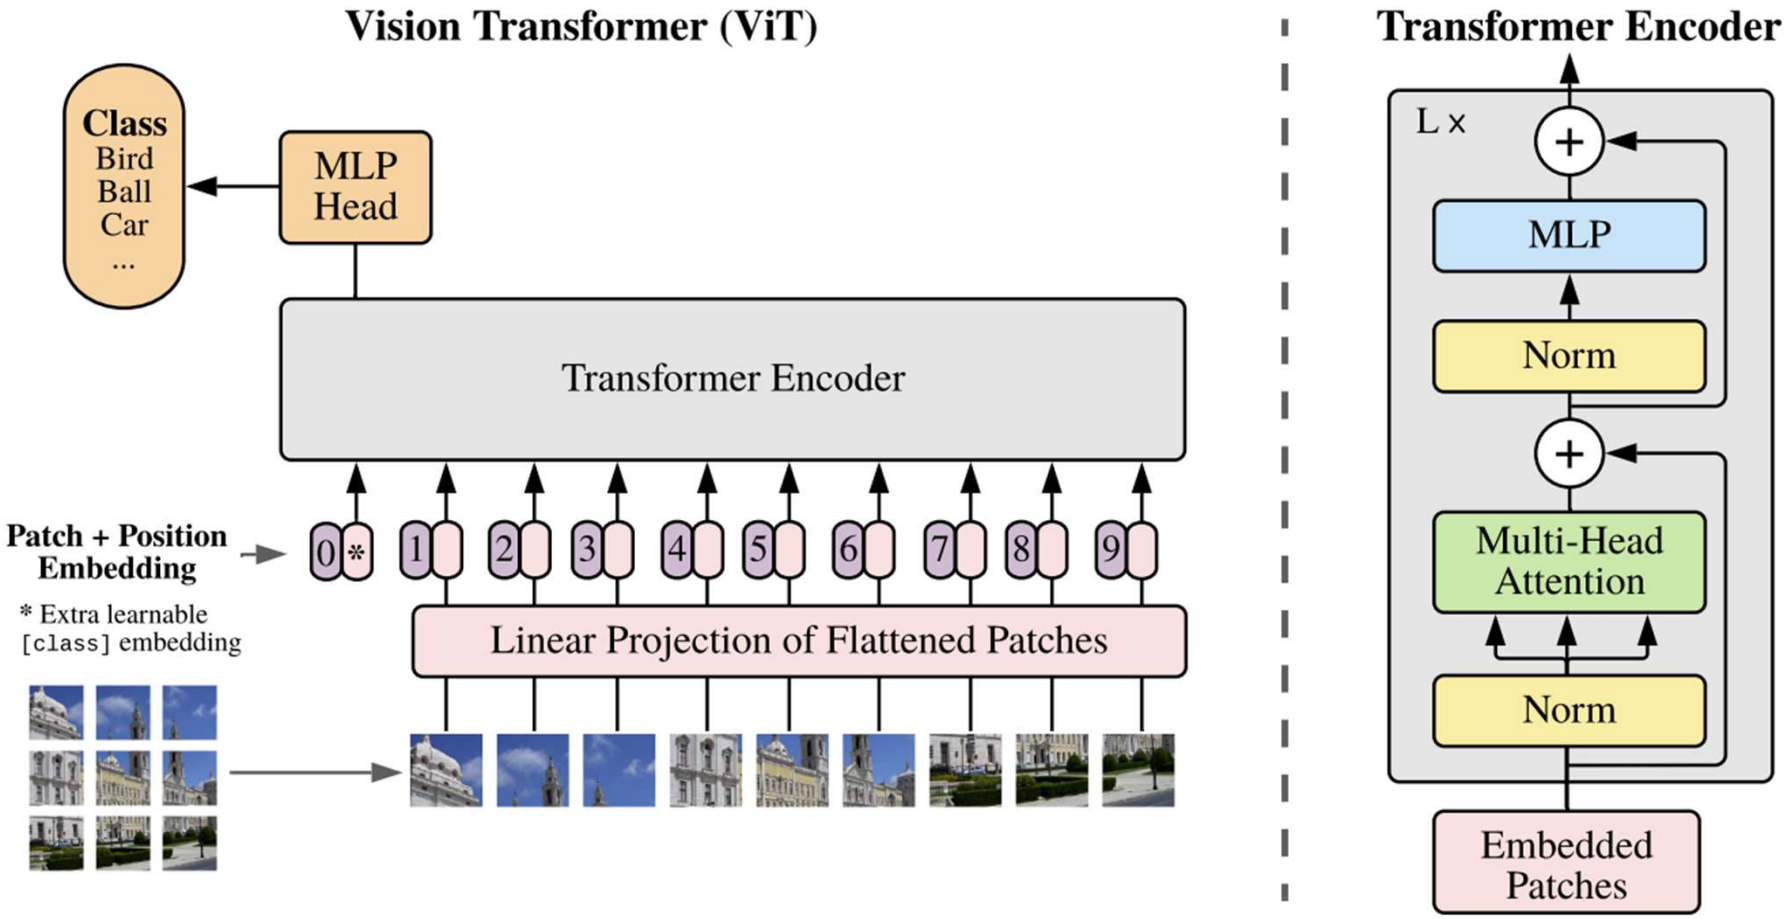

Convolutional neural networks (CNN)

Layer types:

- *Convolutional layer*
- Pooling layer
- Fully-connected layer

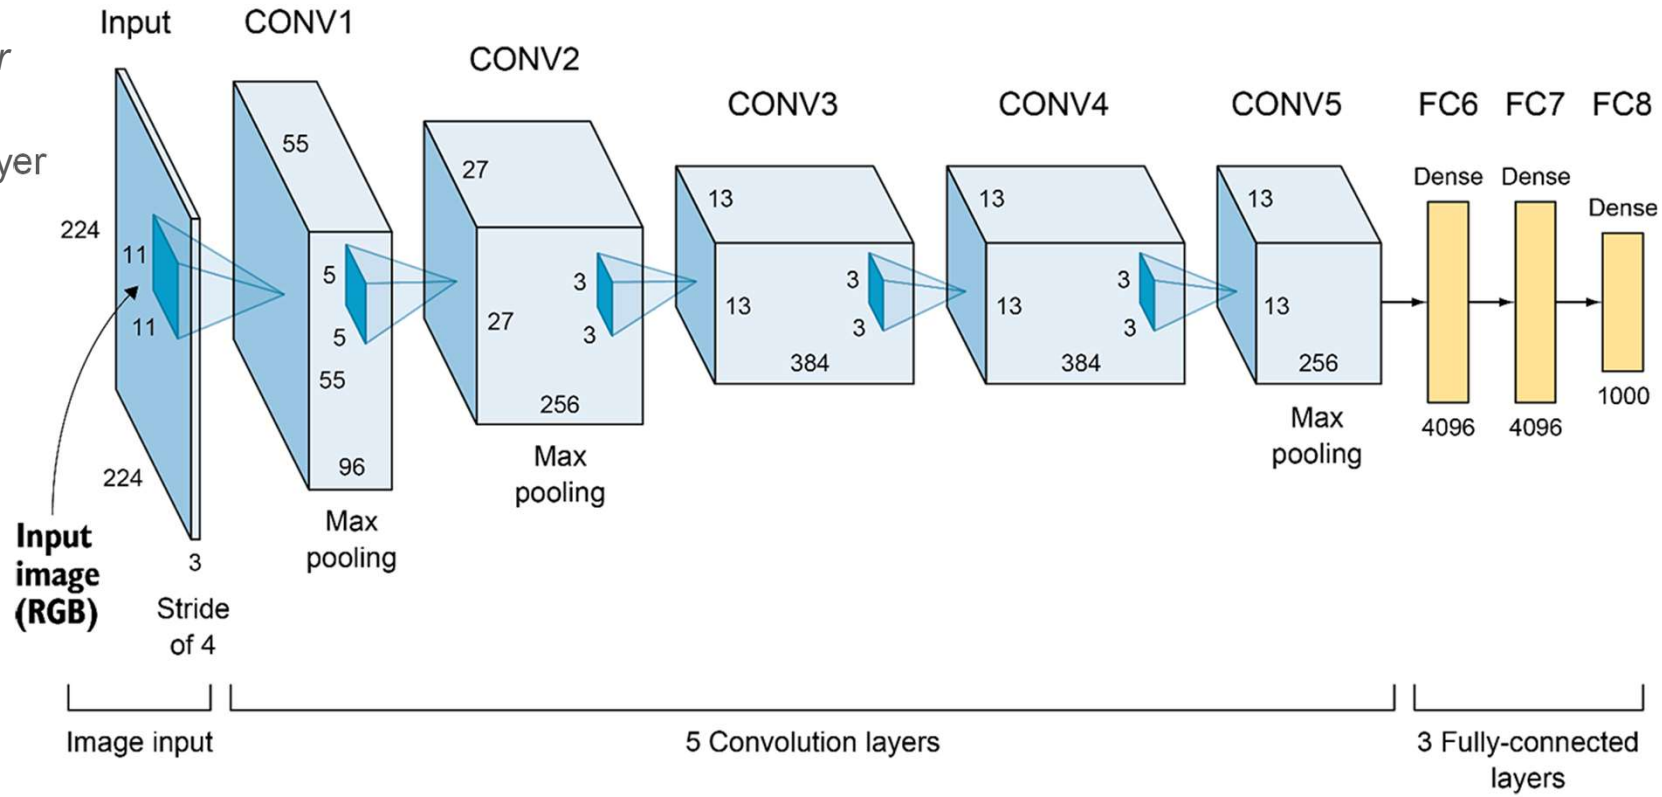

ViT vs CNN for Vision

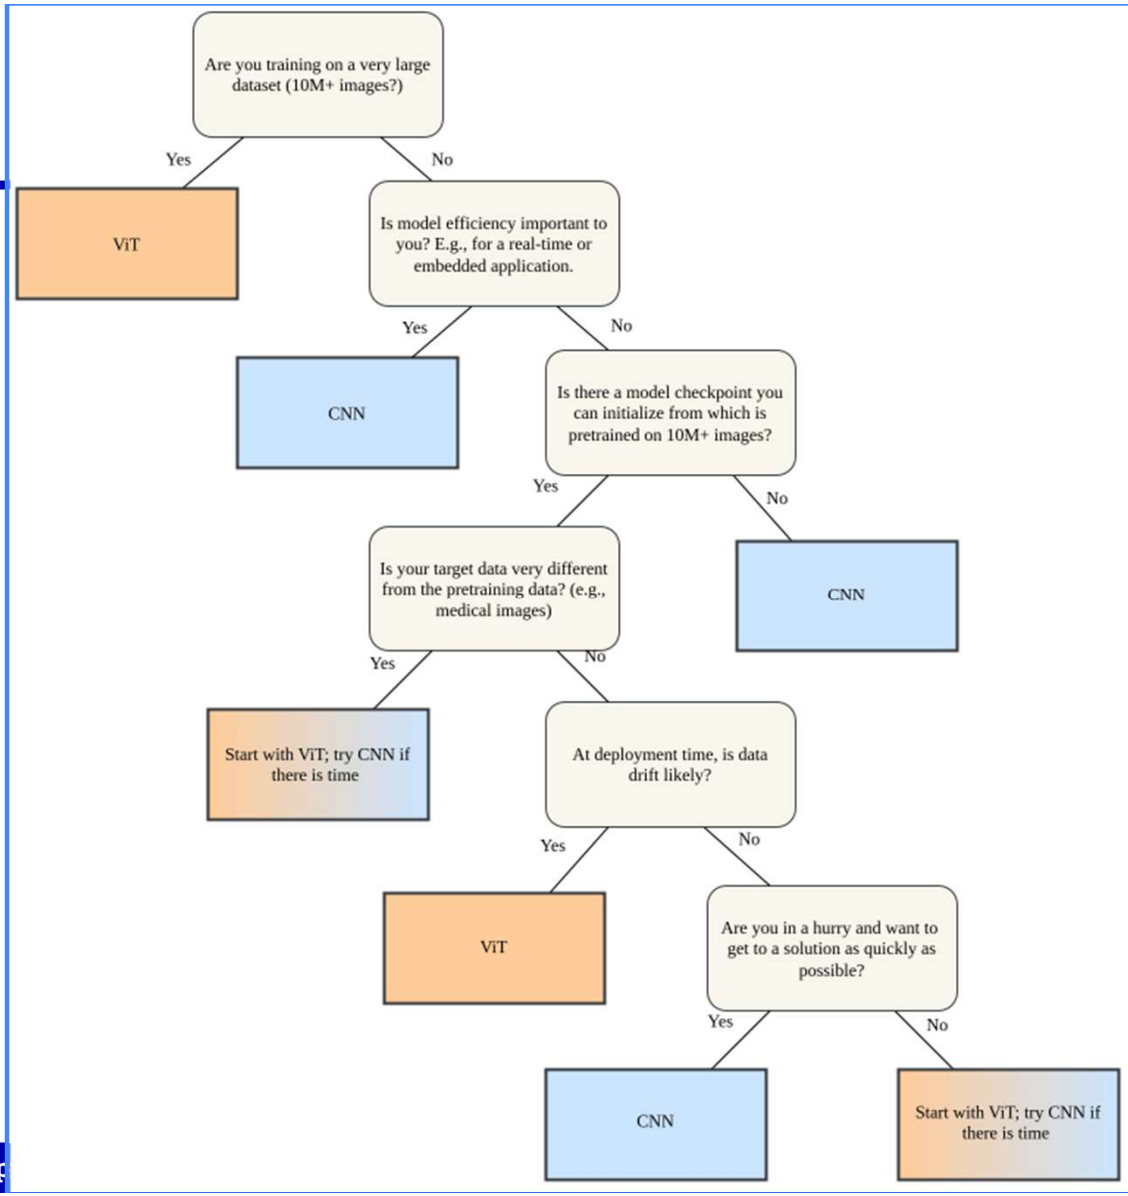

https://tobiasvanderwerff.com/2024/05/15/cnn-vs-vit.html

Further reading

- Attention Is All You Need (https://arxiv.org/pdf/1706.03762)
- The Illustrated Transformer (https://jalammar.github.io/illustrated-transformer/)
- Build a Large Language Model (From Scratch) (https://github.com/rasbt/LLMs-from-scratch)